

Assignment 5.2

author: Rachel Nelson

class: DSC650

Assignment 5.2

Implement the news classifier found in section 3.5 of Deep Learning with Python.

```
20 #3.12 Loading the Reuters dataset
from keras.datasets import reuters
(train_data, train_labels), (test_data, test_labels) = reuters.load_data(num_words=10000)

len(train_data)
len(test_data)

train_data[10]

20 [1,
 245,
 273,
 207,
 156,
 53,
 74,
 160,
 26,
 14,
 46,
 296,
 26,
 39,
 74,
 2979,
 3554,
 14,
 46,
 4689,
 4329,
 86,
 61,
 3499,
 4795,
 14,
 61,
 451,
 4329,
 17,
 12]

21 # 3.13 Decoding newswires back to text

word_index = reuters.get_word_index()
reverse_word_index = dict([(value, key) for (key, value) in word_index.items()])
```

```

decoded_newswire = ' '.join([reverse_word_index.get(i - 3, '?') for i in train_data[0]])

train_labels[10]

21 3

22 # 3.14 Encoding the data
import numpy as np

def vectorize_sequences(sequences, dimension=10000):
    results = np.zeros((len(sequences), dimension))
    for i, sequence in enumerate(sequences):
        results[i, sequence] = 1.
    return results

x_train = vectorize_sequences(train_data)
x_test = vectorize_sequences(test_data)

def to_one_hot(labels, dimension=46):
    results = np.zeros((len(labels), dimension))
    for i, label in enumerate(labels):
        results[i, label] = 1.
    return results

one_hot_train_labels = to_one_hot(train_labels)
one_hot_test_labels = to_one_hot(test_labels)

23 # 3.15 Model definition
from keras import models
from keras import layers

model = models.Sequential()
model.add(layers.Dense(64, activation='relu', input_shape=(10000,)))
model.add(layers.Dense(64, activation='relu'))
model.add(layers.Dense(46, activation='softmax'))

24 # 3.16 Compiling the model
model.compile(optimizer='rmsprop',
              loss='categorical_crossentropy',
              metrics=['accuracy'])

25 # 3.17 Setting aside a validation set
x_val = x_train[:1000]
partial_x_train = x_train[1000:]

y_val = one_hot_train_labels[:1000]
partial_y_train = one_hot_train_labels[1000:]

26 # 3.18 Training the model
history = model.fit(partial_x_train,
                     partial_y_train,
                     epochs=20,
                     batch_size=512,
                     validation_data=(x_val, y_val))

Epoch 1/20
16/16 [=====] - 1s 21ms/step - loss: 2.6140 - accuracy: 0.5167 - val_loss: 1.713
Epoch 2/20

```

```
16/16 [=====] - 0s 13ms/step - loss: 1.4213 - accuracy: 0.7068 - val_loss: 1.285
Epoch 3/20
16/16 [=====] - 0s 13ms/step - loss: 1.0481 - accuracy: 0.7799 - val_loss: 1.140
Epoch 4/20
16/16 [=====] - 0s 13ms/step - loss: 0.8231 - accuracy: 0.8280 - val_loss: 1.024
Epoch 5/20
16/16 [=====] - 0s 13ms/step - loss: 0.6545 - accuracy: 0.8566 - val_loss: 0.961
Epoch 6/20
16/16 [=====] - 0s 14ms/step - loss: 0.5209 - accuracy: 0.8855 - val_loss: 0.956
Epoch 7/20
16/16 [=====] - 0s 13ms/step - loss: 0.4208 - accuracy: 0.9094 - val_loss: 0.913
Epoch 8/20
16/16 [=====] - 0s 13ms/step - loss: 0.3428 - accuracy: 0.9262 - val_loss: 0.892
Epoch 9/20
16/16 [=====] - 0s 13ms/step - loss: 0.2825 - accuracy: 0.9399 - val_loss: 0.884
Epoch 10/20
16/16 [=====] - 0s 13ms/step - loss: 0.2331 - accuracy: 0.9470 - val_loss: 0.889
Epoch 11/20
16/16 [=====] - 0s 14ms/step - loss: 0.2002 - accuracy: 0.9494 - val_loss: 0.939
Epoch 12/20
16/16 [=====] - 0s 13ms/step - loss: 0.1809 - accuracy: 0.9516 - val_loss: 0.924
Epoch 13/20
16/16 [=====] - 0s 13ms/step - loss: 0.1585 - accuracy: 0.9549 - val_loss: 0.964
Epoch 14/20
16/16 [=====] - 0s 13ms/step - loss: 0.1466 - accuracy: 0.9554 - val_loss: 0.989
Epoch 15/20
16/16 [=====] - 0s 14ms/step - loss: 0.1368 - accuracy: 0.9560 - val_loss: 1.018
Epoch 16/20
16/16 [=====] - 0s 13ms/step - loss: 0.1315 - accuracy: 0.9558 - val_loss: 1.061
Epoch 17/20
16/16 [=====] - 0s 12ms/step - loss: 0.1235 - accuracy: 0.9569 - val_loss: 1.067
Epoch 18/20
16/16 [=====] - 0s 13ms/step - loss: 0.1178 - accuracy: 0.9564 - val_loss: 1.023
Epoch 19/20
16/16 [=====] - 0s 13ms/step - loss: 0.1124 - accuracy: 0.9588 - val_loss: 1.244
Epoch 20/20
16/16 [=====] - 0s 13ms/step - loss: 0.1105 - accuracy: 0.9584 - val_loss: 1.060
```


```
27 # 3.19 Plotting the training and validation loss
import matplotlib.pyplot as plt

loss = history.history['loss']
val_loss = history.history['val_loss']

epochs = range(1, len(loss) + 1)

plt.plot(epochs, loss, 'bo', label='Training loss')
plt.plot(epochs, val_loss, 'b', label='Validation loss')
plt.title('Training and validation loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()


plt.show()
```



```
28 # 3.20 Plotting the training and validation accuracy
plt.clf()

acc = history.history['accuracy']
val_acc = history.history['val_accuracy']

plt.plot(epochs, acc, 'bo', label='Training acc')
plt.plot(epochs, val_acc, 'b', label='Validation acc')
plt.title('Training and validation accuracy')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()
plt.show()
```



```
29 # 3.21 Retraining a model from scratch
model = models.Sequential()
model.add(layers.Dense(64, activation='relu', input_shape=(10000,)))
model.add(layers.Dense(64, activation='relu'))
model.add(layers.Dense(46, activation='softmax'))

model.compile(optimizer='rmsprop',
              loss='categorical_crossentropy',
              metrics=['accuracy'])
```



```
batch_size=128,
validation_data=(x_val, y_val))

Epoch 1/20
63/63 [=====] - 1s 8ms/step - loss: 3.2372 - accuracy: 0.3355 - val_loss: 2.6886
Epoch 2/20
63/63 [=====] - 1s 7ms/step - loss: 2.3230 - accuracy: 0.4213 - val_loss: 2.0128
Epoch 3/20
63/63 [=====] - 0s 7ms/step - loss: 1.8037 - accuracy: 0.5471 - val_loss: 1.7330
Epoch 4/20
63/63 [=====] - 0s 7ms/step - loss: 1.6005 - accuracy: 0.5531 - val_loss: 1.6413
Epoch 5/20
63/63 [=====] - 0s 7ms/step - loss: 1.4930 - accuracy: 0.5565 - val_loss: 1.6067
Epoch 6/20
63/63 [=====] - 0s 7ms/step - loss: 1.4118 - accuracy: 0.5877 - val_loss: 1.5854
Epoch 7/20
63/63 [=====] - 0s 7ms/step - loss: 1.3511 - accuracy: 0.5927 - val_loss: 1.5585
Epoch 8/20
63/63 [=====] - 0s 7ms/step - loss: 1.3042 - accuracy: 0.5943 - val_loss: 1.5675
Epoch 9/20
63/63 [=====] - 0s 7ms/step - loss: 1.2641 - accuracy: 0.5957 - val_loss: 1.5886
Epoch 10/20
63/63 [=====] - 0s 7ms/step - loss: 1.2291 - accuracy: 0.5972 - val_loss: 1.5863
Epoch 11/20
63/63 [=====] - 0s 7ms/step - loss: 1.2003 - accuracy: 0.6015 - val_loss: 1.5952
Epoch 12/20
63/63 [=====] - 0s 7ms/step - loss: 1.1731 - accuracy: 0.6061 - val_loss: 1.6302
Epoch 13/20
63/63 [=====] - 0s 7ms/step - loss: 1.1490 - accuracy: 0.6131 - val_loss: 1.6503
Epoch 14/20
63/63 [=====] - 0s 6ms/step - loss: 1.1309 - accuracy: 0.6275 - val_loss: 1.6671
Epoch 15/20
63/63 [=====] - 0s 7ms/step - loss: 1.1094 - accuracy: 0.6353 - val_loss: 1.6717
Epoch 16/20
63/63 [=====] - 0s 7ms/step - loss: 1.0948 - accuracy: 0.6397 - val_loss: 1.7378
Epoch 17/20
63/63 [=====] - 0s 7ms/step - loss: 1.0769 - accuracy: 0.6462 - val_loss: 1.7676
Epoch 18/20
63/63 [=====] - 0s 7ms/step - loss: 1.0639 - accuracy: 0.6646 - val_loss: 1.7743
Epoch 19/20
63/63 [=====] - 0s 7ms/step - loss: 1.0543 - accuracy: 0.6630 - val_loss: 1.8196
Epoch 20/20
63/63 [=====] - 0s 7ms/step - loss: 1.0419 - accuracy: 0.6783 - val_loss: 1.8339
```

31 <keras.callbacks.History at 0x254ffc1c588>

